



SITARE
University

Assignment 1: Training a Simple Neural Network

Course: Artificial Intelligence / Deep Learning

Duration: 1 week

Topic: Deep Neural Networks

Objective

In this assignment, you will implement a very simple Neural Network (NN) with 3 layers (one neuron per layer) and train it to approximate the following function:

$$f(x) = \exp(-\sin(x^2)/2) + x$$

Problem Setup

- Input: x_i (scalar)
- Output: $y_i = f(x_i)$
- Neural Network:
 - Layer 1: 1 neuron (w_1, b_1)
 - Layer 2: 1 neuron (w_2, b_2)
 - Layer 3: 1 neuron (w_3, b_3)
- Activation function: Use ReLU or tanh (your choice, but justify).
- Loss function: Mean Squared Error (MSE).
- Optimizer: Gradient Descent.

Data

Generate 25 samples of x uniformly spaced between $[-10, 10]$.

Compute corresponding $y = f(x)$.

Split randomly into:

- Training set: 20 samples
- Test set: 5 samples

Part A: Mathematical Formulation

1. Forward propagation equations:

$$a[1] = g(w_1x + b_1)$$

$$a[2] = g(w_2a[1] + b_2)$$

$$\hat{y} = g(w_3a[2] + b_3)$$

where $g(\cdot)$ is the chosen activation.

2. Loss function:

$$J = (1/m) * \sum (\hat{y}_i - y_i)^2$$

3. Gradient descent update rule:

$$w_l \leftarrow w_l - \eta \partial J / \partial w_l$$

$$b_l \leftarrow b_l - \eta \partial J / \partial b_l$$

where:

- $l = 1, 2, 3$ (layer index)
- η is the learning rate
- $\partial J / \partial w_l$ and $\partial J / \partial b_l$ are gradients

Part B: Implementation

1. Initialize parameters w_l , b_l randomly.
2. Implement forward propagation.
3. Compute the loss J .
4. Implement backward propagation to compute gradients.
5. Update weights and biases using the rule above.
6. Repeat for multiple epochs (e.g., 1000, 10000 etc. iterations).

Part C: Experiments

1. Train the model on the dataset.
2. Plot:
 - The true function $f(x)$.
 - The predictions \hat{y} after training.
3. Report training and test error.

Part D: Analysis

Answer briefly:

1. Did the simple 3-neuron network approximate the function well?
2. How does the choice of activation (ReLU vs tanh) affect performance?
3. Why might deeper/wider networks perform better on complex functions?

Deliverables

- Jupyter Notebook with code implementation on GitHub.
- Plots of true vs predicted function.
- A short PDF write-up (1–2 pages) with answers to Part D.